


# A1822

09/07/2019

## Autonomous Driving Car in Unity Simulation using Deep Reinforcement Learning



*Figure 1*

Video Link- <https://www.youtube.com/watch?v=og1jGU8T8dE>

### Basic rundown-

RL was used to train a car in Unity simulation to drive a car and not crash for as long as possible. The car is physics based complete with suspension physics/tire physics. The car has constant acceleration at a set rate, speed is capped at 13.3m/s (47.9kmph). Output of the model is a float value between -1 and 1. -1 represents max steering angle to the left while 1 represents max steering to the right.

## Inputs to the model-

- Three raycasts of length 100 units at the front (green lines shown in Fig1 and Fig-2), returns distance detected.
- Eight raycasts of length 20 units (black lines shown in Fig-1 and Fig 2), returns distance detected.
- Speed of the car.
- Drift value, which is calculated via dot product of car velocity direction and car forward direction. Measures if the car is oversteering/sliding or stable.
- Angle of the wheel.
- Velocity of the car is x-axis, y-axis and z axis.




Figure 2

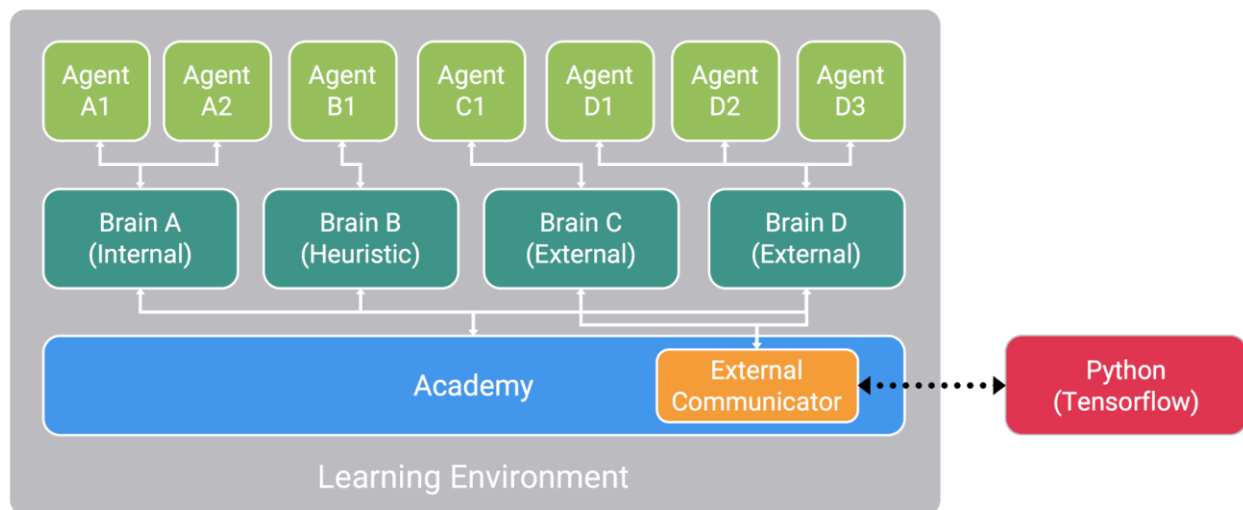


Figure 3 Unity ML-Agents

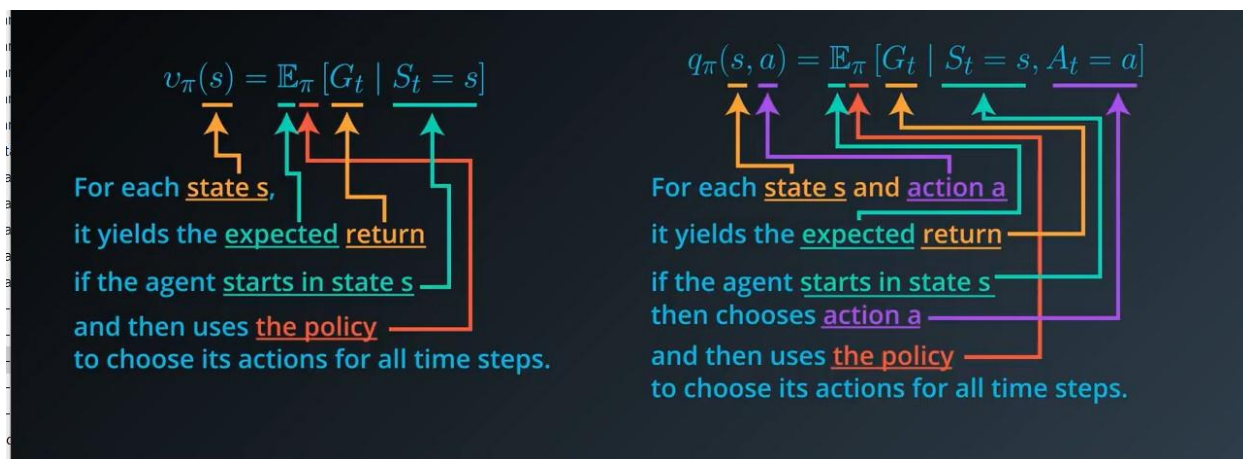


Figure 4




Figure 5

## Rewards and Punishments-

### Rewards-

- (forward ray distance/6000f)
- (forward left ray distance/10000f)
- (forward right ray distance/10000f)
- if(driftValue>0.90) then 0.01 reward
- if(car\_Speed>5) then 0.001 reward
- if(car\_Speed>10) then 0.1 reward

### Punishments-

- $-(100f - \text{forward ray distance})/50000f$
- $-(100f - \text{forward left ray distance})/90000f$
- $-(100f - \text{forward right ray distance})/90000f$
- if any forward green rays less than 4f then -0.08 reward
- if(driftValue<0.90) then -0.05 reward
- if any forward green rays less than 1f then resets reward to -1 for the current step and resets car position

## Training info-

Car agent was trained for 1.8 million “steps”. Car agent has 10000 steps to not crash, after 10000 or a crash, car agent respawns. Inputs, raycast angles/position, car agent spawn points were tweaked as necessary between steps as model was routinely saved. Trainer used was Proximal Policy Optimization (PPO).

### Hyperparameters-

```
trainer: ppo  
  
beta: 5.0e-3  
  
epsilon: 0.2
```

```

lambda: 0.95
learning_rate: 3.0e-4

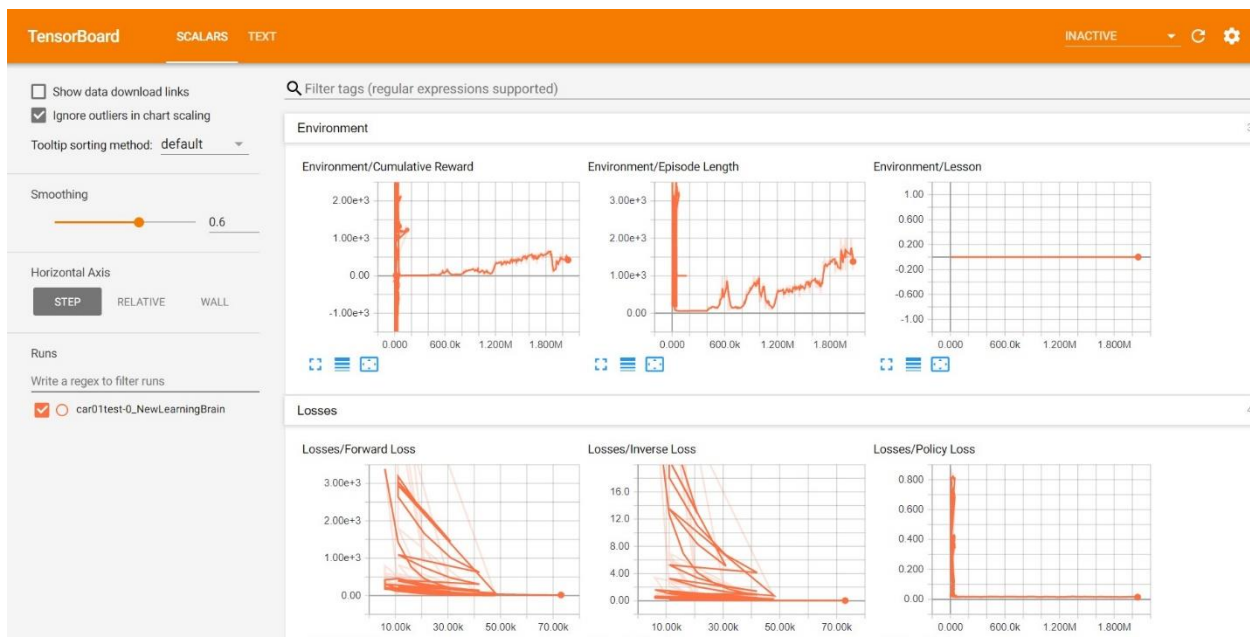
memory_size: 256

sequence_length: 64


use_recurrent: false
use_curiosity: false
curiosity_strength: 0.01
curiosity_enc_size: 128
normalize: true
num_epoch: 3
time_horizon: 1000
batch_size: 2024
buffer_size: 20240
gamma: 0.995
max_steps: 5e6
summary_freq: 3000
num_layers: 3
hidden_units: 512

```

Tensorboard-



Losses/Value Loss



## Overfitting in RL-

Overfitting is “to fall into a local minimum”. In this sim the RL models generates its own data. An agent looks for an optimal strategy, interacting with the environment. If environment reacts “kindly”, the strategy is reinforced, the agent is “motivated” to repeat the same actions over and over aiming to get higher reward. If there is an exploit in the rewards, the agent inevitably learns to exploit it. To avoid this the rewards and punishments were tweaked until the agent learned to generalize.

## Video link-

<https://www.youtube.com/watch?v=og1jGU8T8dE>

## Project executable link-

<https://drive.google.com/open?id=16Vgichx38PfQy13Udjuevkn3HS5ajYRY>

## Assets Used-

- Unity Vehicle Pack
- Mountain Road

## Conclusion-

In RL the most difficult thing to achieve is the right balance of reward function. Doing this project resulted in a better understanding of that. This project maybe further improved by addition of things such as tire pressure, gravity measurements, virtual cameras, traffic, off-road environments, pedestrians, giving the agent acceleration control, braking control and so on.